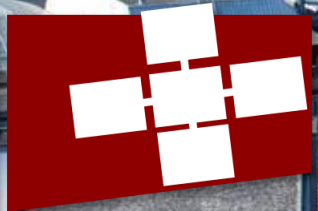
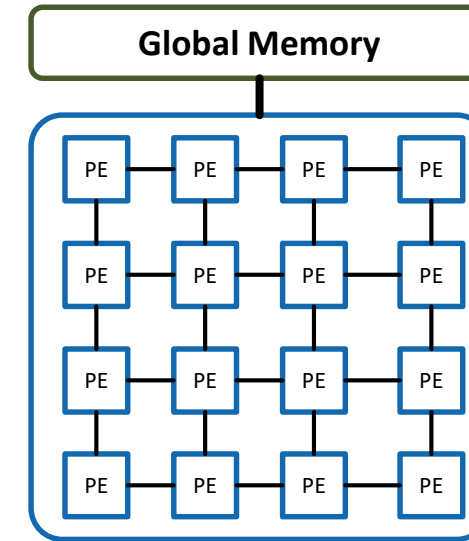



# ASA: Scheduling






# Scheduling





In traditional task scheduling, a task can start only when all the predecessors completed (compute-then-communicate)



# Streaming Scheduling

We want to enable **streaming communications** between tasks




**In this way we exploit spatial (pipeline) parallelism and reduce off-chip memory accesses**

Solving this adds complexity to an already complicated problem:

- We need to understand whether it is better to stream or not, by building **streaming blocks**
- We need to understand how to schedule these streaming blocks

# Input and assumptions



**Input:** a graph  $G = (V, E)$ , where:

- The nodes represent tasks (operations) in which the computation can be decomposed and that can be executed on a PE
- Edges represent data movements and dependencies. The labels indicate the number of data elements being transferred

**Output:** a schedule for  $G$  that minimizes its running time (makespan)

**Target architecture:**

- a spatial device composed by homogenous PEs
- PEs are fully connected
- Backing memory. PEs can always communicate through main memory

**Communication:** completely decoupled (solid edges) or pipelined (dashed)

**Assumptions:**

- Blocking read semantics
- Tasks are not preemptible
- Co-scheduling: weak or strict
- Communications occur w/o contention

# Two subproblems

We need to solve two inter-related problems




**Streaming:** decide which edge implements a pipelined communication and which not

**Scheduling:** schedule the tasks considering pipelined communication

We start by addressing the Scheduling problem

# Scheduling

Let's assume that we have the DAG with the type of communications




---

## Algorithm 1. List-Scheduling

---

```


while there are tasks to be scheduled do
  Identify a highest priority task  $n$  (e.g., from a list);
  Choose a processor  $p$  for  $n$ ;
  Schedule  $n$  on  $p$  at  $est(n, p)$ ;
end
  
```

---

Candidates: HEFT, PEFT, ... (something specific for homog. Computation TBD)

- take into account streaming (a children can start before its parent finishes)
- best-effort: streaming tasks are co-scheduled if it can do it, and it is useful, otherwise not
- Backpressure may be a problem

# Tasks



The running time is given by the volume of data being ingested/consumed

$$T_t = \max\{N, M\}$$

**Alternatively:** we need additional input data to tell us how long does it takes to compute a task, what is its initiation interval, what is the latency ...

A node may be slow down by a slower child. This affects its running time and the running time of the other children

The backpressure effect must be taken into account while scheduling the tasks

**Alternatively:** assume that we can always scale a task. This must be taken into account in the analysis

# TODO

- Find good heuristics and use them as reference
- Understand how to model backpressure
- Play with some toy example
- [Streaming]